

Predictive Phase Continuation from Solvable Models toward Realistic Learning Systems

Yuma Ichikawa

July 20, 2026

Abstract

We introduce *predictive phase continuation*: effective variables are calibrated at solvable nodes and used to predict, without refitting, phase boundaries at held-out deformations. The test is deliberately stronger than curve similarity. It requires semantic retention of the order parameter, largest-size prediction against continuous, first-order-like, and crossover alternatives, and a preregistered intervention near the predicted critical window. Our controlled four-domain benchmark separates outer disorder from inner stochasticity and concentrates additional seeds only near inferred boundaries. Across 46080 registered tasks, the median held-out boundary error falls from 0.011 to 0.008 after adding a domain-specific macrovariable, a 26.1% improvement. Critical-window allocation removes 57.1% of uniform compute while preserving the registered quality target. Transformer and diffusion are the deep controlled cases; reinforcement learning and multi-agent dynamics test the same predictive grammar without assuming common equilibrium exponents. An artifact audit also identifies why saturated legacy order parameters cannot support boundary claims. The present evidence establishes a controlled benchmark, not yet external validity on natural-data or open-weight endpoints.

1 Introduction

Solvable learning models increasingly exhibit sharp changes in recovery, attention mechanism, memorization, policy support, and collective order. The central unresolved question is not whether such transitions can be drawn, but whether a theory calibrated before realism is added predicts an unseen boundary afterward. We ask:

Can a small set of macroscopic variables calibrated in a solvable model predict phase boundaries in a held-out learning system, and can prediction failure identify the missing variable?


This question changes phase continuation from a descriptive atlas into a falsifiable transport problem. For an edge $u \rightarrow v$ in an assumption graph, we fit effective parameters only at u , predict $\hat{g}_{c,v}$, and score the prediction against an untouched confirmatory grid at v . A failed bridge is scientifically useful only when an augmented observable restores prediction on the same holdout.

Our four primary claims are: predictive boundary transport, localization of non-additive assumption interactions, recovery by a new macrovariable, and compute-efficient intervention inside the predicted critical window. The large experiment registry, immutable artifacts, and plotting grammar are supporting infrastructure rather than scientific claims.

2 Closest prior work and the unresolved gap

Solvable attention already predicts a positional-semantic transition in a tied low-rank layer [5], while attention-indexed models extend closed-form predictions to full-width matrices [3]. Thus observing a transition or a spectral outlier is not our novelty; the test is transport to a hidden

Predictive phase continuation



Every claim: discovery -> confirmation -> hidden holdout

Figure 1: Predictive phase continuation. A solvable anchor calibrates effective variables; all boundary and intervention scores are evaluated on a blinded node.

Domain	Known result	Assumption removed here	Hidden prediction
Attention	positional-semantic transition	tied/low-rank surrogate	deformed g_c
Diffusion	speciation and collapse times	oracle/global score	locality-shifted time
RL	entropy and Goodhart dynamics	exact verifier	gold boundary
Agents	field-driven crossover	homogeneous mean field	coupling tipping point

Table 1: Existing results are treated as anchors, not rediscovered claims.

deformation. Diffusion theory separates speciation from memorized-point collapse [2], and recent DiT experiments report near-concurrent symmetry breaking and nonlocality onset [9]. We therefore predict relative critical times rather than rediscovering these diagnostics.

In RL, policy-entropy collapse has a mechanistic account and targeted controls already exist [4]; proxy optimization can degrade gold quality as KL budget grows [8]. Entropy alone is consequently not an order parameter. In multi-agent systems, global-flip experiments show that apparent consensus can be field-driven rather than interaction-driven [6]. Effective-team-size laws and matched coordination controls further show that nominal agent count is not independent evidence [1, 7].

3 Predictive phase continuation

Let $\mathcal{O}_{\text{cal}}$ denote registered calibration observables and let $\mathcal{O}_{\text{solv}}(\theta)$ be a solvable surrogate. We estimate

$$\hat{\theta}_{\text{eff}} = \arg \min_{\theta} D\left(\mathcal{O}_{\text{cal}}^{\text{target}}, \mathcal{O}_{\text{solv}}(\theta)\right), \quad (1)$$

then predict $\hat{g}_c$, the order curve, and the response to an intervention. The bridge error is evaluated only on withheld quantities,

$$\mathcal{E}_{\text{bridge}} = D\left(\mathcal{O}_{\text{holdout}}^{\text{target}}, \mathcal{O}_{\text{solv}}(\hat{\theta}_{\text{eff}})\right). \quad (2)$$

For each edge we also record semantic retention $S_{\text{sem}} = I(m; Z)/H(Z)$ and normalized predictive error $|g_c - \hat{g}_c|/(\sigma_{g_c} + \epsilon)$.

Two assumptions need not combine additively. Their interaction is

$$\Delta_{ij}g_c = g_c^{ij} - g_c^i - g_c^j + g_c^0. \quad (3)$$

The operational outcome vocabulary includes stable and renormalized continuation, splitting, merging, rounding, new-phase insertion, computational–statistical separation, path dependence, and semantic failure. These labels are not inferred from visual smoothness.

We distinguish macrostate perplexity $K_{\text{eff}}^{\text{macro}} = \exp H(Z)$ from correlation-based independent components and spectral effective rank. Each PhaseCard stores estimator, probe distribution, normalization, units, finite-size variable, intervention, and null model.

4 Statistical protocol

Discovery, confirmation, and holdout are separate datasets. The largest variant is absent from every fit. The first confirmatory stage uses 16 outer disorder seeds and four inner stochastic replicates. After boundary localization, 16 additional outer seeds are added only to the adjacent three control values. Inner replicates never count as independent disorder draws. Intervals use a hierarchical bootstrap over outer-seed means, while within- disorder variance is reported separately.

For every transition, six system sizes are evaluated. A continuous width $w_N \propto N^{-1/2}$, a first-order-like width $w_N \propto N^{-1}$, and a smooth crossover $w_N = w_\infty + aN^{-1/2}$ compete by prediction of the largest size. Exponents are called asymptotic only if window, minimum-size, correction-term, held-out-size, and scaling-relation checks all agree; otherwise they are finite-size effective exponents.

The vector of system sizes is never collapsed by convention: representation dimension, reverse-time resolution, MDP horizon, and agent population have domain-specific meanings. Likewise, Binder ratios are used only where their symmetry assumptions are satisfied. RL uses occupancy and strategy participation; multi-agent analysis separates intrinsic field from social coupling with a global flip.

All empirical figures display raw outer-seed points beneath means and 95% intervals. One panel has one physical control. Cross-domain comparison is performed only after domain-specific normalization.

5 Artifact audit and preregistration

Before designing new experiments, we audited the preceding 49,440-run atlas. The audit records metric minima, maxima, rounded unique counts, saturation at zero or one, and the fraction of zero-width intervals for each domain–family–metric group. It flags 71 groups as dominated by saturation. In particular, the former representative table used an order parameter equal to one with zero uncertainty in several learned and cross-domain families. That table is removed rather than reinterpreted. Non-saturated signed order, task error, and response diagnostics remain available as supplementary checks.

The new registry contains four primary claims, one primary order per claim, explicit nulls, an untouched holdout variant, and a missingness rule that forbids incomplete aggregation. Artifact records retain all inner values, outer seeds, device, configuration, and task identifiers. Cluster-specific paths remain outside the repository.




Figure 2: Anchor calibration with hierarchical 95% intervals. The lower panel exposes finite-size and disorder residuals rather than hiding them behind a theory line.

6 Transformer and diffusion deep cases

6.1 Attention recovery

The controlled attention problem retains positional and semantic latent signals. The primary order is signed recovery rather than a clipped accuracy; secondary observables are semantic retention, precursor strength, specialization entropy, and correlation-based component fraction. Six sizes and a fixed confirmatory control grid determine the boundary. The holdout adds a non-additive deformation that is absent from calibration.

6.2 Diffusion dynamics

The diffusion case separates semantic speciation, nonlocality, memorization, and trajectory commitment. Reverse time is not reused as a control for score noise or locality; the secondary axis modifies locality before normalized boundary comparison. The central prediction is the holdout speciation boundary, with precursor nonlocality as the augmenting macrovariable.

7 Driven validation: reinforcement and multi-agent systems

The RL benchmark separates oracle occupancy overlap, strategy entropy, gold reward, proxy reward, Goodhart gap, and effective rollout fraction. The control plane is KL pressure versus verifier noise. Productive exploration, specialization, entropy collapse, and reward hacking therefore cannot be collapsed into a single entropy curve.




Figure 3: Transformer holdout. Raw outer seeds, hierarchical intervals, and all six finite sizes are shown on the domain-specific signal control.

The multi-agent benchmark uses truth-conditioned consensus and explicitly varies intrinsic field h and social coupling J . Positive and negative fields implement the global-flip control. Effective component fraction and error reproduction are reported alongside nominal population size, preventing shared bias from masquerading as independent consensus.

8 Held-out prediction, breakdown, and intervention

The base surrogate attains median boundary error 0.011, whereas the augmented surrogate attains 0.008, improving held-out transport by 26.1%. The mean anchor residual is 0.002, and the largest registered non-additive shift is 0.048. These quantities come directly from the immutable aggregate.

Prediction failure is tested rather than explained post hoc: the same holdout is scored by base and augmented surrogates. Critical-window interventions allocate the expensive operation only within a fixed interval around the predicted boundary and compare against always-on, always-off, random-window, and matched-compute controls. The registered allocation removes on average 57.1% of uniform compute while retaining 87.3% of the always-on quality gain and exceeding a random matched-compute window by 0.019.

9 Limitations

The new experiment is a controlled Tier-B stochastic benchmark. It does not establish natural-language, natural-image, open RLVR, or open-weight LLM-agent external validity. The previous repository contains trainable Tier-B+ smoke models, but their saturated primary order failed the




Figure 4: Diffusion holdout. Speciation is shown across six sizes; memorization and nonlocality are distinct registered secondary diagnostics.

Domain	Base error	Augmented error	Compute saving
Diffusion	0.013	0.010	57.1%
Multiagent	0.011	0.008	57.1%
Reinforcement	0.005	0.005	57.1%
Transformer	0.011	0.010	57.1%

Table 2: Automatically generated held-out results. No saturated legacy order parameter enters this table.

present audit and is not promoted to evidence. External endpoints require version-pinned models, natural-data holdouts, matched token/FLOP controls, and independent implementations.

Our finite sizes can distinguish competing predictive trends but do not by themselves prove an asymptotic universality class. RL and agent systems are driven and potentially nonstationary; thermodynamic language is used only for operational dynamic regimes. Critical-window quality is benchmark-specific, and compute savings do not transfer automatically to wall-clock savings on a different architecture.

Finally, an adaptive design can bias inference if the first-stage boundary is used to redefine the estimand. We avoid this by freezing the original grid, adding seeds rather than moving controls, and retaining both stages in the manifest.




Figure 5: RL holdout curves with hierarchical 95% intervals. Gold reward (top) and the proxy-gold Goodhart gap (bottom) separate productive scaling from reward overoptimization at three verifier-noise levels.

10 Conclusion

Predictive phase continuation asks whether a solvable model forecasts an unseen boundary, not whether two curves look alike. The framework makes failure local and useful: non-additive assumption interactions identify where transport breaks, and an additional macrovariable must recover prediction on the same blinded system. Nested disorder, competing transition models, adaptive boundary replication, and critical-window interventions convert a large atlas into four falsifiable claims. The controlled benchmark is a step toward, rather than a substitute for, natural-data external validation.

A PhaseCard and reproducibility contract

Every condition records eight taxonomy axes: thermodynamic level, symmetry, disorder, dynamics, system-size vector, control variable, order family, and evidence class. Outcome labels include stable or renormalized continuation, splitting, merging, rounding, new-phase insertion, computational-statistical separation, path dependence, and semantic failure.

Production uses only a portable Spark profile. The script rejects any partition not beginning with `spark_`; repository files contain no site path. A run is immutable after its atomic result write. Aggregation fails on a missing task. The complete JSON stores all outer means, inner variance, bootstrap interval, boundary estimate, competing width fits, interaction shifts, and intervention summaries. Figures are generated only from that JSON. Every figure has physical size 6.4 by 4.8 inches, inward ticks, 12-point base type, larger axis labels, and preregistered color, line, and




Figure 6: Multi-agent holdout curves with hierarchical 95% intervals. Positive and negative intrinsic fields provide the global-flip control needed to separate field-driven alignment from interaction-driven order.

marker order.

B Legacy atlas status

The 58-family atlas remains a searchable diagnostic benchmark and code-coverage target. It is not evidence for the four claims above. Programs A–N, secondary architectures, and common-coordinate diagnostics move to the supplement because they otherwise obscure the held-out prediction experiment.

References

- [1] Blaž Bertalančič and Carolina Fortuna. The ringelmann effect in multi-agent llm systems: A scaling law for effective team size. *arXiv preprint arXiv:2606.02646*, 2026.
- [2] Giulio Biroli, Tony Bonnaire, Valentin de Bortoli, and Marc Mézard. Dynamical regimes of diffusion models. *Nature Communications*, 15:9957, 2024.
- [3] Fabrizio Boncoraglio, Emanuele Troiani, Vittorio Erba, and Lenka Zdeborová. Bayes optimal learning of attention-indexed models. In *Advances in Neural Information Processing Systems*, 2025. arXiv:2506.01582, revised 2026.
- [4] Ganqu Cui, Yuchen Zhang, Jiacheng Chen, et al. The entropy mechanism of reinforcement learning for reasoning language models. *arXiv preprint arXiv:2505.22617*, 2025.




Figure 7: Predicted versus observed boundaries for the blinded variant. Marker denotes domain and horizontal bars are hierarchical 95% intervals on the observed boundary.

- [5] Hugo Cui, Freya Behrens, Florent Krzakala, and Lenka Zdeborová. A phase transition between positional and semantic learning in a solvable model of dot-product attention. In *Advances in Neural Information Processing Systems*, 2024.
- [6] Cristiano De Nobili. Collective alignment in llm multi-agent systems: Disentangling bias from cooperation via statistical physics. *arXiv preprint arXiv:2605.10528*, 2026.
- [7] Yubin Kim, Ken Gu, Chanwoo Park, et al. Towards a science of scaling agent systems. *arXiv preprint arXiv:2512.08296*, 2026.
- [8] Rafael Rafailov, Yaswanth Chittapu, Ryan Park, Harshit Sikchi, Joey Hejna, Chelsea Finn, et al. Scaling laws for reward model overoptimization in direct alignment algorithms. *arXiv preprint arXiv:2406.02900*, 2024.
- [9] Yifan F. Zhang, Fangjun Hu, Guangkuo Liu, Mert Okayay, and Xun Gao. Concurrence of symmetry breaking and nonlocality phase transitions in diffusion models. *arXiv preprint arXiv:2605.04830*, 2026.




Figure 8: Held-out error before and after the registered macrovariable is added. Error bars summarize variation over size and secondary control, not pooled inner replicates.